

How to Quickly Deploy Machine Learning Model Anywhere

Dr. Erekle Magradze

Ilia State University
Computer Science Program



Challenge

- We moved from the age of discovery to the age of implementation.
- Many great startups launched in the garages – obviously having a garage is not enough :)
- It's not only enough to have a nice Jupyter notebook to demonstrate the results – the expectation is to have the MVP (Minimal Viable Product) with the ML model, as soon as possible – that's your competitive advantage;

How to do that???

Problems

You have trained the model – now you'd like to show it to your customer over the ocean – asking to do the following steps


- Install Python/go/Julia/Rust/R
- Install all the necessary packages and libraries
- Setup environment variables (don't forget to open this and that ports)
- Copy my code to your machine – don't forget to split them in several files
- Run the code with required parameters

Eventually

**THIS IS
NOT
WORKING**

Perception vs Reality

Perception



Google Cloud



Credit: Hidden Technical Debt of Machine Learning Systems, D. Sculley, et al.

@dsculley/hello

Reality



Figure 1: Only a small fraction of real-world ML systems is composed of the ML code, as shown by the small black box in the middle. The required surrounding infrastructure is vast and complex.

Google Cloud

Credit: Hidden Technical Debt of Machine Learning Systems, D. Sculley, et al.

@dsculley/hello

I am a Software or ML engineer ...

And I don't care about the technical stuff required for the deployment - well, nowadays it's not like that anymore ;)

- Use containerization technology – in particular the docker <https://www.docker.com/> use community edition, it's free
- Once installed and it's running you can run the following command in your terminal (in case of Mac and Linux) or in your powershell (in case of windows)

Here is the command

docker run --rm -p 5000:5000 datascienceexplorer/classifier

After running this command go to the following url

<http://localhost:5000/apidocs/>

- This is a ready application with Machine Learning model in it – based on IRIS dataset, you can try it over the swagger i.e. the link shown above

OKay, but what exactly is docker?

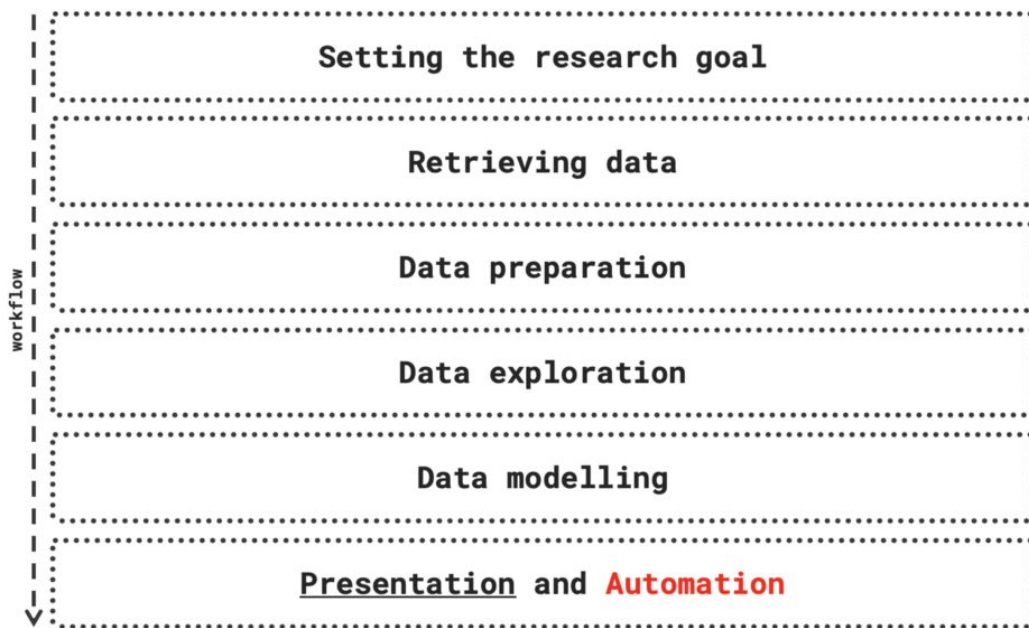
Docker is your friend :)

- Docker is a containerization tool, working and running on majority of operating systems;
- It helps to pack all the necessary packages and dependencies in one file – called container image and deploy it on any OS without installation of additional software or opening the ports and doing sysadmin or system integrator tasks – it's all there for free, out of the box
- How it's possible because of containerization concept

Containers? Long story - short



Steps of ML project



It's great when you have lots of time to wait for ANN or CNN model training - automate it and run it, that's fine!

When you need to give a quick response - we need to keep the model in memory

In case of python use pickle - object serialization package

```
1 import pickle
2
3 with open('./model.pkl', 'wb') as model_pkl:
4     pickle.dump(knn, model_pkl)
```


Let's have a look to the Jupyter notebook

<https://github.com/hermag/docker-session/blob/master/SupervisedMachineLearningClassification.ipynb>

- More or less all should be clear if you are doing ML or Data Science
- The most important part for us is this

Save Model

```
In [52]: import pickle
```

```
In [55]: with open('./model.pkl', 'wb') as model_pkl:  
         pickle.dump(knn, model_pkl)
```

How to use pickled model

We can write simple code to load the pickled model and use it




```
1  import pickle
2  # Import all the packages you need for your model below
3  import numpy as np
4  from sklearn.neighbors import KNeighborsClassifier
5
6  # Load the model into memory
7  with open('./model.pkl', 'rb') as model_pkl:
8      knn = pickle.load(model_pkl)
9
10 # Unseen data (create a new observation for testing)
11 unseen = np.array([[3.2, 1.1, 1.5, 2.1]])
12 result = knn.predict(unseen)
13
14 # Print result to the console
15 print('Predicted result for observation ' + str(unseen) + ' is: ' + str(result))
```

But you might face the following problem
To solve this we need to install the
missing package

```
Traceback (most recent call last):
  File "main.py", line 4, in <module>
    from sklearn.neighbors import KNeighborsClassifier
ImportError: No module named sklearn.neighbors
```

This is not enough - we need a service

REST API Design



In case of python we can use Flask

```
1 import pickle
2 # Import all the packages you need for your model below
3 import numpy as np
4 import sys
5 from sklearn.neighbors import KNeighborsClassifier
6 # Import Flask for creating API
7 from flask import Flask, request
8
9 # Load the trained model from current directory
10 with open('./model.pkl', 'rb') as model_pkl:
11     knn = pickle.load(model_pkl)
12
13 # Initialise a Flask app
14 app = Flask(__name__)
15
16 # Create an API endpoint
17 @app.route('/predict')
18 def predict_iris():
19     # Read all necessary request parameters
20     sl = request.args.get('sl')
21     sw = request.args.get('sw')
22     pl = request.args.get('pl')
23     pw = request.args.get('pw')
24
25     # Use the predict method of the model to
26     # get the prediction for unseen data
27     unseen = np.array([[sl, sw, pl, pw]])
28     result = knn.predict(unseen)
29
30     # return the result back
31     return 'Predicted result for observation ' + str(unseen) + ' is: ' + str(result)
32
33
34 if __name__ == '__main__':
35     app.run()
```

Finally the dockerfile

To wrap the flask application we finally need the dockerfile to build the docker image

```
1  FROM python:3.7
2  WORKDIR /app
3  EXPOSE 5000
4
5  COPY ./requirements.txt .
6  RUN pip install -r requirements.txt
7
8  COPY . .
9  CMD python main.py
```

How our folder structure should look?

```
.  
├── Dockerfile  
├── README.md  
├── SupervisedMachineLearningClassification.ipynb  
├── main.py  
├── model.pkl  
└── requirements.txt
```

To build the container image: ***docker build -t mymlmodel***

To start the container using the built image: ***docker run -name contname -d mymlmodel***

To expose the container ports: ***docker run -name contname -p 5000:5000 -d mymlmodel***

To attach the storage: ***docker run -name contname -v localdrive:/container/path -p 5000:5000 -d mymlmodel***

All materials are available here

<https://github.com/hermag/docker-session>